

Goroutines and Channels in Go

Session 15

Golang course by Exadel

31 Nov 2022

Sergio Kovtunenکو

Lead backend developer, Exadel

Agenda

- ▶ Goroutines
- ▶ Channels: bi-directional and uni-directional, buffered and unbuffered
- ▶ Channels Axioms
- ▶ Select statement
- ▶ Blocking vs. non-blocking flows
- ▶ Rules for using channels
- ▶ Next time...

Before everything else...

▶ You can write production-grade applications in Go without knowing much about:

- goroutines
- channels
- mutexes, etc.

▶ We need goroutines to **better utilize** our hardware!

Goroutines

Goroutines 🧨

- ▶ The go statement launches a function call as a goroutine

```
go f()  
go f(x, y, ...)
```

- ▶ A goroutine runs **concurrently** (but not necessarily in parallel)

- ▶ A goroutine is a **thread of control within the program**, with its **own local variables and stack**.

- ▶ We have no goroutine **ID** ! No "handler" associated with goroutine!

- ▶ *Much cheaper to create and schedule* than operating system threads.

- ▶ **Main** goroutine vs others.


- ▶ Example: `code/01_goroutines_test.go`

Source: "Introduction to Go" by Dave Chaney

(<https://github.com/davecheney/introduction-to-go>)

Understanding Concurrent vs. Parallel Execution

- ▶ Concurrency is about dealing with lots of things at once
- ▶ Parallelism is about doing lots of things at once
- ▶ Concurrency is about structure, parallelism is about execution



Source: <https://talks.golang.org/2012/waza.slide#8>

(<https://talks.golang.org/2012/waza.slide#8>)

Goroutines recap

▶ Goroutines are user-space **lightweight** threads.

- Go runtime **scheduler** will manage them, not the OS. The scheduler is of **work-stealing** type.

▶ Very **lightweight**!

- Owns a **small stack** for local variables: around 2k per 1 goroutine + **stack can grow** as needed
- Faster to create and do context switch
- You can have millions of them in your application

▶ Goroutines are multiplexed onto OS threads

- We can limit the number of OS threads for program written in Go using GOMAXPROCS env var
- `runtime.GOMAXPROCS(32)`

`go` statement (1/2)

- ✓ A "go" statement starts the execution of a function call as an independent control, or **goroutine**, within the same address space.
- ✓ The expression must be a function or method call; it cannot be a variable.
 - Calls of built-in functions are restricted as for expressions.
- ✓ The function value and parameters are evaluated as usual in the program with a regular call, program execution does not wait for the function to return.
 - Instead, the function begins executing independently in a new goroutine.
 - When the function terminates, its goroutine also terminates.
 - If the function has any return values, they are discarded.
- ✓ Example:

```
go Server()  
go func(ch chan<- bool) { for { sleep(10); ch <- true }} (c)
```


`go` statement (2/2)

- ✓ You don't need "goroutine ID" at all.
 - In Go there is NO goroutine exposed.
 - Potential problems include:
 - Because of the high potential for abuse, there is no identifier for the current goroutine in Go.
 - This may seem draconian, but this actually preserves package ecosystem: *it does not matter if you start a*

// That is, for any method of function F:

`F()`

// is almost exactly equivalent to:

`done := make(chan struct{})`

`go func() {`

`defer close(done)`

`F()`

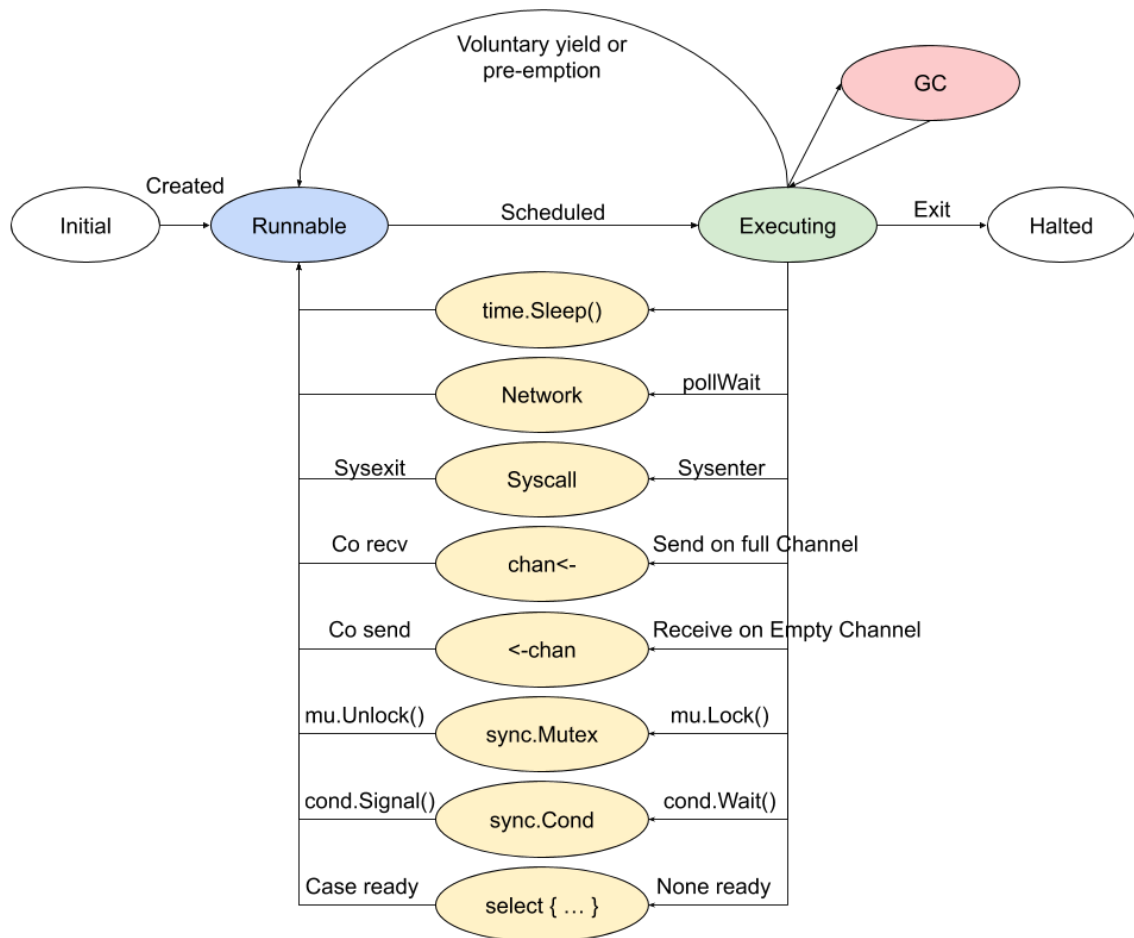
`}`

`<-done`

- ✓ When you spawn goroutines, make it clear when - or whether - th
 - Try to keep concurrent code simple enough that goroutine .
 - If that just isn't feasible, document when and why the gor

Goroutine lifecycle

Goroutine Lifecycle



Source: "cmd/trace: problems and proposed improvements #33322" (<https://github.com/golang/go/issues/33322>)

Goroutine context switching

▶ To understand context switching there is terminology:
M/G/P :

▶ **G** - represents a user space **Goroutine**

- it has its own stack and blocking info

▶ **P** - represents a logical entity for available **Processors** or a scheduling context.

- The number of P is pre-decided (GOMAXPROCS) and *fixed during the run.*

▶ **M** - stands for **Machines**, they are a representation for OS threads (like POSIX thread)

- **M** will run a set of **G's**
- Every **Machine (M)** needs a **P** to run **G**.

▶ The same names M/G/P is in use in internal Go's source code.

Understand Golang scheduler

▶ Awesome 30-mins video: [GopherCon 2018: Kavya Joshi - The Scheduler Saga](https://www.youtube.com/watch?v=YHRO5WQGh0k) (<https://www.youtube.com/watch?v=YHRO5WQGh0k>)

▶ Another awesome video: [Dmitry Vyukov — Go scheduler: Implementing language with lightweight concurrency](https://www.youtube.com/watch?v=K11rY57K7k)

(<https://www.youtube.com/watch?v=K11rY57K7k>)

▶ Watch them! :)

▶ Also useful article: [Go Runtime Scheduler Design Internals](https://freethreads.wordpress.com/2019/01/23/go-runtime-scheduler-design-internals/)

(<https://freethreads.wordpress.com/2019/01/23/go-runtime-scheduler-design-internals/>)

Channels

Channels: basic info

- ✓ A `channel` provides a mechanism for concurrently executing functions and receiving values of a specified element type.
 - The value of an uninitialized channel is `nil`.
 - A `nil` channel is never ready for communication.
- ✓ A single channel may be used in:
 - send statements OR
 - receive operations OR
 - calls to the built-in functions `cap()` and `len()` by any number for further synchronization.
 - Channels act as first-in-first-out queues.
 - For example, if one goroutine sends values on a channel and another receives them, the values are received in the order they were sent.

Channels: channel direction

- ✓ The optional `<-` operator specifies the **channel direction**, *send*
 - If *no direction* is given, the channel is **bidirectional**.
 - A channel may be *constrained* only to send or only to rece

```
chan T           // can be used to send and receive values of type T
chan<- float64   // can only be used to send float64s
<-chan int       // can only be used to receive ints
```

- ✓ The `<-` operator associates with the leftmost `chan` possible:

```
chan<- chan int   // same as chan<- (chan int)
chan<- <-chan int // same as chan<- (<-chan int)
<-chan <-chan int // same as <-chan (<-chan int)
chan (<-chan int)
```


- ✓ A new, *initialized channel* value can be made using the built-in the channel type and an **optional capacity** as arguments:

```
unbufferedChannel := make(chan int)
bufferedChannel := make(chan int, 100)
```

Channels: buffered or non-buffered

- ✓ Channels can be **buffered** or **non-buffered**.
- ✓ Unbuffered channels.
 - If *channel capacity* is zero or absent, the channel is unbuf
 - The communication succeeds only when both a *sender* and *rece*
- ✓ Buffered channels.
 - Channels can be buffered. Provide the buffer length as the initialize a buffered channel.
 - Sends to a buffered channel *block* only when the buffer is f
 - Receives *block* when the buffer is empty.
- ✓ A **channel** may be *closed* with the built-in function **close()**.
 - The multi-valued assignment form of the receive operator *re* was sent before the channel was closed.


Visualization of non-buffered channels



Source: "stackoverflow: What are channels used for?"

(<https://stackoverflow.com/a/39831976>)

Visualization of buffered channels



Source: "stackoverflow: What are channels used for?"

(<https://stackoverflow.com/a/39831976>)

Closing a channel

- ✓ For a channel `c`, the built-in function `close(c)` records that no more values will be sent on the channel.
 - It is an error if `c` is a *receive-only channel*.
- ✓ Sending to or closing a *closed* channel causes a run-time panic.
- ✓ Closing the `nil` channel also causes a run-time panic.
- ✓ It is completely normal to not to close channels manually.
 - They will be automatically garbage collected once related function finishes.

Best Practice: Formal arguments for goroutine functions

- ✓ Prefer using formal arguments for the channels you pass to goroutines in global scope.
 - You can get *more compiler checking* this way, and *better modularity*
 - Example:

```
func main() {
    c := make(chan string)

    for i := 1; i <= 2; i++ {
        go func(i int, co chan<- string) {
            for j := 1; j <= 2; j++ {
                co <- fmt.Sprintf("hi from %d.%d", i, j)
            }
        }(i, c)
    }

    for i := 1; i <= 4; i++ {
        fmt.Println(<-c)
    }
}

// OUTPUT:
// Hi from 2.1
// Hi from 2.2
// Hi from 1.1
// Hi from 1.2
```

Source: "Multiple goroutines listening on one channel"
StackOverflow ([http://stackoverflow.com/questions/15715605/multiple-goroutines-listening-on-one-](http://stackoverflow.com/questions/15715605/multiple-goroutines-listening-on-one-channel/15721380#15721380)

[channel/15721380#15721380](http://stackoverflow.com/questions/15715605/multiple-goroutines-listening-on-one-channel/15721380#15721380))

Best Practice: Avoid both reading and writing on the same channel

- ✓ **Avoid both reading and writing on the same channel** in a particular **'main'** one).
 - Otherwise, deadlock is a much greater risk.
- ✓ It is generally a good principle to *view buffering as a performance*
 - If your program does not deadlock without buffers, it won't (but the converse is not always true).
 - So, as another rule of thumb, start without buffering then

Channel Axioms: send to a nil channel blocks forever

- ✓ A send to a **nil** channel blocks forever:

```
package main

func main() {
    var c chan string
    c <- "let's get started" // deadlock
}
```

Source: "Channel Axioms" by Dave Cheney

(<https://dave.cheney.net/2014/03/19/channel-axioms>)

Channel Axioms: receive from a nil channel blocks forever

- ✓ A receive from a `nil` channel blocks forever:

```
package main
import "fmt"

func main() {
    var c chan string
    fmt.Println(<-c) // deadlock
}
```

- *Reason why that happened:*
 - If the channel is not initialised then its *buffer size*
 - If the size of the channel's buffer is zero, then th
 - If the channel is unbuffered, then a send will block ready to receive.
 - If the channel is `nil` then the sender and receiver h they are both blocked waiting on independent channel

Source: "Channel Axioms" by Dave Cheney

(<https://dave.cheney.net/2014/03/19/channel-axioms>)

Channel Axioms: send to a closed channel panics

- ✓ A send to a **closed channel** *panics*:

```
package main
import "fmt"

func main() {
    var c = make(chan int, 100)
    for i := 0; i < 10; i++ {
        go func() {
            for j := 0; j < 10; j++ {
                c <- j
            }
            close(c)
        }()
    }
    for i := range c {
        fmt.Println(i)
    }
}
```

Source: "Channel Axioms" by Dave Cheney

(<https://dave.cheney.net/2014/03/19/channel-axioms>)

Channel Axioms: receive from a closed channel returns the zero value immediately

- ✓ A receive from a **closed channel** returns the *zero value immediately*

```
package main
import "fmt"

func main() {
    c := make(chan int, 3)
    c <- 1
    c <- 2
    c <- 3
    close(c)
    for i := 0; i < 4; i++ {
        fmt.Printf("%d ", <-c) // prints 1 2 3 0
    }
}
```

Source: "Channel Axioms" by Dave Cheney

(<https://dave.cheney.net/2014/03/19/channel-axioms>)

Next time...

We will continue next time

▶ We didn't discuss today:

- select statement
- goroutine gotcha's

Homework:

▶ "3 Golang Channel Red Flags" (<https://hackmongo.com/post/three-golang-channel-red-flags/>)

▶ Must read: Channel Axioms (<https://dave.cheney.net/2014/03/19/channel-axioms>)

▶ Channels in Go (<https://go101.org/article/channel.html>)

▶ Principles of designing Go APIs with channels

(<https://inconshreveable.com/07-08-2014/principles-of-designing-go-apis-with-channels/>)

▶ Understanding real-world concurrency bugs in Go

(<https://blog.acolyer.org/2019/05/17/understanding-real-world-concurrency-bugs-in-go/>)

▶ Just for information, not to be used: A pattern for overcoming non-determinism of Golang select statement

(<https://medium.com/@pedram.esmaeeli/a-pattern-for-overcoming-non-determinism-of-golang-select-statement-139dbe93db98>)

Next session...

 Session16:

Most useful packages: context, sync

- Context package
- Sync package

Thank you

Golang course by Exadel

31 Nov 2022

Sergio Kovtunenکو

Lead backend developer, Exadel

skovtunenکو@exadel.com (mailto:skovtunenکو@exadel.com)

<https://github.com/skovtunenکو> (https://github.com/skovtunenکو)

[@realSKovtunenکو](http://twitter.com/realSKovtunenکو) (http://twitter.com/realSKovtunenکو)

